

# Variational Autoencoder for Facial Image Generation with Bayesian Optimization

Probabilistic Machine Learning

---

Annalisa Paladino



# Introduction

Dataset: lego facial images



Goal: to achieve high-quality image generation




How?

- **VAE** to generate facial image
- **Bayesian optimization** to optimize VAE's hyperparameters


# Variational Autoencoder (VAE)

**ENCODER:** takes as input an image and compress it into a compact, **latent** representation. The VAE outputs the parameters that define a *probability distribution* (mean and log-variance) over the latent space

**DECODER:** samples a point from the distribution in the latent space and attempts to **reconstruct** the original input as accurately as possible



# Variational Autoencoder (VAE)



The input  $x$  is mapped through the encoder into a latent space  $z$  from which we can reconstruct  $\hat{x}$  via the decoder

Goal:  $\hat{x}$  is as **similar** as possible to  $x$

$$p(x, z|\theta) = p(x|z, \theta)p(z)$$

$$q(z|x, \phi) = N(\mu_z(x, \phi), \sigma_z^2(x, \phi) \cdot I)$$

# Reparametrization trick

Decompose a random variable  $z$  from the distribution into a **deterministic component** and a **stochastic component**


1. **Introduce a new independent random variable  $\varepsilon$ :**  $\varepsilon$  serves as a source of randomness entirely separate from the sampling process within the VAE
2. **Reparameterize the latent variable:** create a scaled version of the new random variable ( $\varepsilon$ ) using the standard deviation predicted by the encoder

$$z = \mu + \sigma \cdot \varepsilon$$

deterministic outputs  
from the encoder


source of randomness

$$\text{Loss} = \text{MSE} + \text{KLD}$$



### Mean Squared Error

Measures the **reconstruction error** between the original image and the reconstructed image



### Kullback-Leibler Divergence

Measures the **difference** between the **learned latent distribution** and a **standard normal** distribution, encouraging the model to learn a regularized latent space

# Bayesian Optimization

Goal: find the set of input parameters  $x$  that **maximize** a function  $f(x)$

$$\operatorname{argmax}_{x \in A} f(x)$$

In our case:

- The input of  $f$  is a vector of hyperparameters, from a subset of  $\mathbb{N}^2$
- The output is the validation loss  $\mathbb{R}^+$

# Bayesian Optimization - *properties*

- $A$  is a set of points  $x$ , which rely upon less (or equal to) than 20 dimensions ( $\mathbb{R}^d, d \leq 20$ )
- The objective function  $f$  is **continuous**
- $f$  is **expensive** to evaluate due to its computational cost



# Bayesian Optimization - *properties*

- $f$  is a **black-box** : it lacks structure like concavity or linearity
- $f$  is **not differentiable**: we observe only  $f(x)$  and no first- or second- order derivatives
- Our focus is on finding **global optimum** → the goal is to have a guide when exploring the hyperparameter's space

# Bayesian Optimization - *components*

**Surrogate model:** a statistical model to **approximate** the objective function  $f(x)$

**Acquisition function:** function that **guides** where to sample next

# Algorithm

Place a Gaussian process prior on  $f$

Observe  $f$  at  $n_0$  points according to an initial space-filling experimental design. Set  $n = n_0$

**While**  $n \leq N$  **do**

    Update the posterior probability distribution on  $f$  using all available data

    Let  $x_n$  be a maximizer of the acquisition function over  $x$ , where the acquisition function is computed using the current posterior distribution

    Observe  $y_n = f(x_n)$

    Increment  $n$

**End while**

Return the point evaluated with the largest  $f(x)$

# ***Surrogate model - Gaussian Process***

A **Gaussian Process** (GP) is a stochastic process indexed by a continuous variable  $x \in \mathbb{R}^n$ , such that all finite dimensional distributions are multivariate Gaussian

GP is uniquely defined by:

- $\mu: \mathbb{R}^n \rightarrow \mathbb{R}$  function that computes the **mean** at every point  $x$
- $K: \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$  function that computes the **covariance** between any two points

# ***Surrogate model - Gaussian Process***

$$f \sim GP(\mu, K) \Leftrightarrow \forall x_1, \dots, x_n \in \mathbb{R}^n, p(f) = p(f(x_1), \dots, f(x_N)) = N(\underline{\mu}, K)$$

With:

$$\underline{\mu} = (\mu(x_1), \dots, \mu(x_N))$$

$$K = (K_{ij}) \text{ with } K_{ij} = k(x_i, x_j)$$

Let's fix the mean to 0 and the kernel to be the Matern with parameters

$$\rho = 1, \quad \nu = 2.5$$



# ***Acquisition function* - Expected Improvement**

**Improvement:** how much better the function value  $f(x)$  at a given point  $x$  is compared to the best known function value  $f_n^*$  achieved at step  $n$

$$[f(x) - f_n^*]^+$$

Where

$$a^+ = \max(a, 0)$$

$$f_n^* = \max_{m \leq n} f(x_m)$$

But...  $f(x)$  is unknown ➡ **EXPECTED IMPROVEMENT**

# *Acquisition function* - Expected Improvement

Given a surrogate model (Gaussian Process) that provides a **mean** and **standard deviation** of the predicted function value at point  $x$ , the **expected improvement** is

$$\text{EI}_n = \mathbb{E}[(f(x) - f_n^*)^+]$$

Where

$$a^+ = \max(a, 0)$$

$$f_n^* = \max_{m \leq n} f(x_m)$$

# Hyperparameters optimized


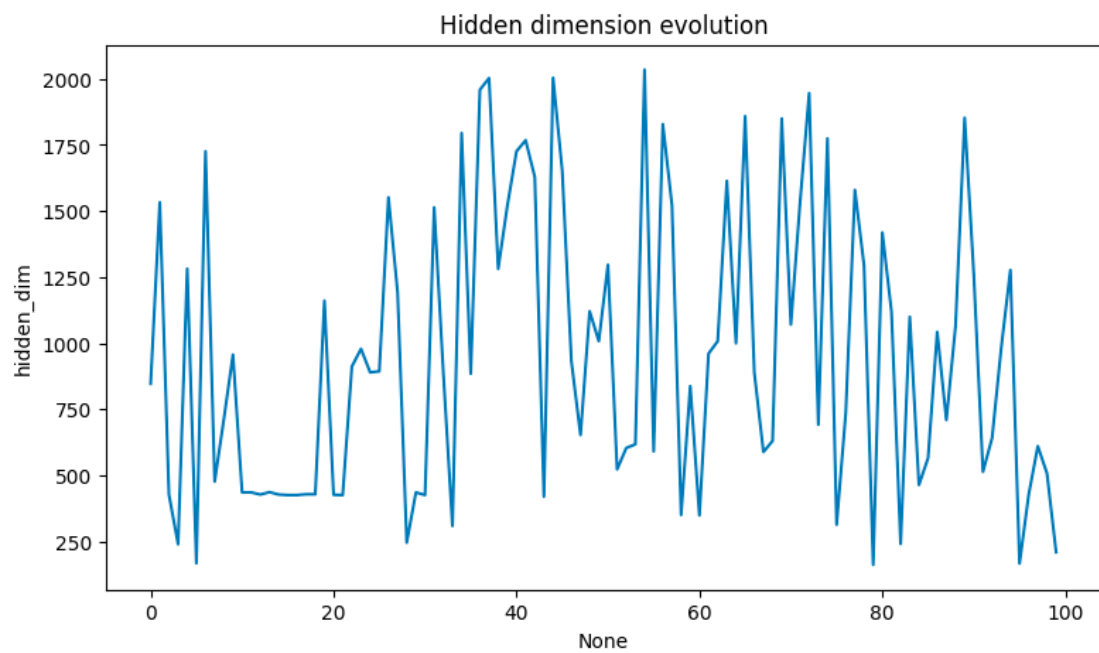
- **Hidden dimension:** Controls the number of neurons in the fully connected layers of the model
- **Latent dimension:** Controls the size of the latent space, which impacts the model's ability to generate varied outputs

Bayesian Optimization samples the hyperparameters over:

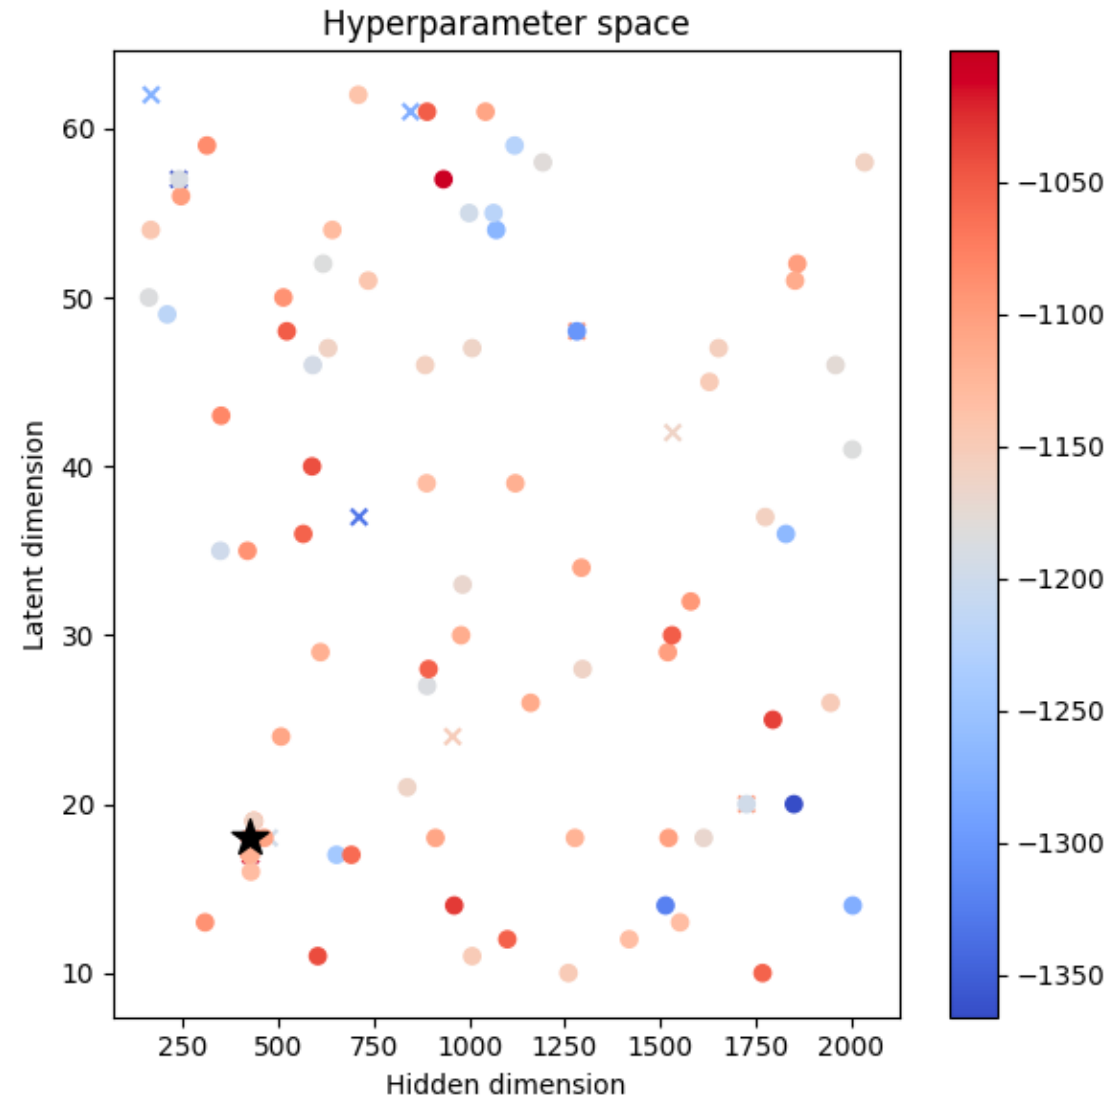
- 10 random points to explore the hyperparameter space
- 90 further iterations to refine the search around promising hyperparameters

Epochs: 50

# Hyperparameters evolution



# Hyperparameter space





# Experimental set-up

The optimization process identified the best hyperparameters:

**Hidden dimension:** Optimal size found within the range of 128 to 2048

**Latent dimension:** Optimal size found between 8 and 64

The VAE model is instantiated with these values!

# Experimental set-up

Optimizer: Adam

Epochs: 400

Initial learning rate:  $1e-3$

Final learning rate:  $5e-7$


Learning rate scheduling: to decay the learning rate over time

Gradient clipping: to avoid instability and to prevent exploding gradients


clip\_grad\_value: 1.0

Batch size: 64

# Loss




# Image Generation



# Face Morphing

The latent space is **continuous** and **smooth**



We can **interpolate** between points in this space to generate smoothly transitioning variations of our data





# Conclusions

The model learned to encode images into a latent space and generate reconstructions

Bayesian optimization tuned the model's hyperparameters, improving the model's performance without requiring extensive manual tuning

The model demonstrated smooth transitions between facial images through interpolation



**Thank you for your  
attention !**